

Lab 6: Generative Models

Yu-Hsiang, Chen - 313553004
August 20, 2024

1 Introduction

In this lab, we delve into the implementation of a Conditional Denoising Diffusion Probabilistic Model (DDPM) to generate synthetic images conditioned on specified multi-label conditions. The objective is to create a model that, given a set of conditions such as "red sphere," "yellow cube," and "gray cylinder," can generate images that accurately reflect these attributes. To evaluate the model's performance, the generated images will be passed through a pre-trained classifier, and their accuracy will be compared against provided testing datasets.

The DDPM is fundamentally a generative model that builds on the principles of diffusion processes. Mathematically, the DDPM involves a forward process where noise is incrementally added to an image over a series of T time steps, resulting in a noisy image x_T . This process is modeled as a Markov chain, where each step t introduces Gaussian noise to the image:

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{\alpha_t}x_{t-1}, (1 - \alpha_t)\mathbf{I})$$

Here, α_t is a noise schedule parameter, and \mathbf{I} is the identity matrix. The reverse process, which is the core of the DDPM, seeks to iteratively remove the noise from x_T to generate a clear image x_0 . This reverse process is parameterized by a neural network $p_\theta(x_{t-1} | x_t, y)$, where y represents the conditioning labels:

$$p_\theta(x_{t-1} | x_t, y) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t, y), \Sigma_\theta(x_t, t, y))$$

The mean μ_θ and covariance Σ_θ are learned parameters, and the network θ is trained to predict the original image x_0 from a noisy version x_t while respecting the conditioning information y .

In our implementation, we have introduced three key components:

- **Data Loader:** A custom data loader that systematically pairs images with their respective multi-label conditions, ensuring proper dataset preparation for training and testing phases.
- **Model Architecture:** We designed a Conditional DDPM with three attention mechanisms. The attention layers are incorporated to enhance the model's ability to focus on specific features that correspond to the given conditions y . The attention mechanism can be described mathematically as a weighting function that assigns different importance to different parts of the image representation, allowing the model to prioritize relevant features during the denoising process.
- **Noise Schedule and Sampling Method:** The noise schedule α_t was carefully designed to balance the trade-off between the diffusion (adding noise) and denoising phases.

We employed a custom sampling method to improve the accuracy and efficiency of generating synthetic images.

Our expectations for this experiment are twofold: first, to generate high-quality synthetic images that match the specified conditions; second, to achieve high classification accuracy when these images are evaluated using a pre-trained ResNet18 classifier.

2 Implementation Details (Describe how you implement your model, including your choice of DDPM, noise schedule)

2.1 Data Loader Implementation

The data loader is a crucial component of our implementation, responsible for preparing the dataset for both training and testing. We designed the data loader to handle image files and their corresponding multi-label annotations. The implementation is encapsulated in the `SyntheticObjectsDataset` and `SyntheticObjectsLabelsOnlyDataset` classes, which are tailored to handle both images with labels and labels only, respectively.

The `SyntheticObjectsDataset` class takes as input the directory containing images, a JSON file with label annotations, and an object dictionary that maps object names to indices. The images are loaded using the PIL library, transformed to the appropriate size (64x64 pixels), normalized, and converted to tensors. The labels are encoded as one-hot vectors, where each label is represented by a vector of zeros with a single one at the index corresponding to the label.

The data loader implementation is shown below:

```
class SyntheticObjectsDataset(Dataset):
    def __init__(self, images_dir, labels_file, object_dict, transform=None):

        self.images_dir = images_dir
        self.transform = transform
        self.object_dict = object_dict

        # Load the label file
        with open(labels_file, 'r') as f:
            self.labels = json.load(f)

        # List of image file names
        self.image_files = list(self.labels.keys())

    def __len__(self):
        return len(self.image_files)

    def __getitem__(self, idx):

        img_name = self.image_files[idx]
        img_path = os.path.join(self.images_dir, img_name)

        # Load the image
        image = Image.open(img_path).convert('RGB')

        # Apply transformations
```

```

if self.transform:
    image = self.transform(image)

# Get the labels and encode them
label_names = self.labels[img_name]
labels = torch.zeros(len(self.object_dict))

for label in label_names:
    label_index = self.object_dict[label]
    labels[label_index] = 1.0

return image, labels

```

Listing 1: SyntheticObjectsDataset class

For scenarios where only labels are needed, such as during testing, the SyntheticObjectsLabelsOnlyDataset class is used. This class reads the labels from the JSON file and encodes them similarly, but without loading any images.

The data loader is initialized using the `get_data_loader` function, which returns a `DataLoader` object that can be used directly in training and evaluation loops. This function also handles the choice between returning both images and labels or labels only, based on the `return_labels_only` parameter.

2.2 Model Design and UNet Architecture

The core of our model is a Conditional Denoising Diffusion Probabilistic Model (DDPM), which leverages a U-Net architecture for generating high-quality synthetic images based on specific label conditions. The U-Net model is particularly well-suited for image generation tasks because of its ability to capture both local and global features through its symmetrical down-sampling and up-sampling paths.

In our implementation, we used the `UNet2DModel` from the `huggingface diffusers` library as the backbone for our DDPM. The model is designed to process 64x64 pixel images, with three input channels (RGB) and three output channels (also RGB). The architecture includes six blocks in both the down-sampling and up-sampling paths, with attention layers strategically placed in the middle blocks to enhance the model's ability to focus on relevant features related to the input labels.

The model is defined as follows:

```

class CustomDDPM(nn.Module):
    def __init__(self, num_classes=24, model_dim=512):
        super(CustomDDPM, self).__init__()

        self.embedding_dim = model_dim // 4
        # we can change for different embedding for labels
        self.embedding_layer = nn.Linear(num_classes,
                                         model_dim)

        self.unet_model = UNet2DModel(
            sample_size=64,
            in_channels=3,
            out_channels=3,
            layers_per_block=2,
            block_out_channels=[
                self.embedding_dim,
                self.embedding_dim,
                self.embedding_dim * 2,
                self.embedding_dim * 2,
                self.embedding_dim * 4,
                self.embedding_dim * 4
            ],

```

```

        down_block_types=[
            "DownBlock2D",
            "DownBlock2D",
            "DownBlock2D",
            "AttnDownBlock2D",
            "AttnDownBlock2D",
            "AttnDownBlock2D"
        ],
        up_block_types=[
            "AttnUpBlock2D",
            "AttnUpBlock2D",
            "AttnUpBlock2D",
            "UpBlock2D",
            "UpBlock2D",
            "UpBlock2D"
        ],
        class_embed_type="identity"
    )

    def forward(self, input_tensor, time_step, labels):
        embedded_class = self.embedding_layer(labels)
        output = self.unet_model(input_tensor, time_step,
                                embedded_class)
        return output.sample

```

Listing 2: CustomDDPM class with U-Net architecture

The U-Net model's structure allows it to effectively capture multi-scale features, which are crucial for generating detailed and accurate synthetic images. The attention mechanisms within the middle blocks of the U-Net allow the model to focus on specific parts of the image that are most relevant to the conditioning labels, thus improving the quality and accuracy of the generated outputs. We experimented with different attention configurations and found that incorporating more attention layers in the latter stages of the model led to better performance in terms of image quality.

2.3 Training Implementation

Training the DDPM involves optimizing the model to minimize the difference between the generated noise and the actual noise added during the diffusion process. We implemented the training loop using the PyTorch framework, where we defined the loss function as Mean Squared Error (MSE) between the predicted noise and the ground truth noise, also we'll use Adam for optimizer.

Our training configuration includes the following parameters:

- **Epochs:** 200
- **Batch Size:** 32
- **Learning Rate:** 0.0002
- **Timesteps:** 1000
- **Noise Scheduler:** A linear noise scheduler implemented using `DDPMScheduler` from the `diffusers` library

The training loop is as follows:

```

def train_model(model, train_loader, test_loader, device,
               noise_scheduler):
    model.to(device)
    criterion = nn.MSELoss()
    optimizer = optim.Adam(model.parameters(), lr=config['learning_rate'])

```

```

writer = SummaryWriter(config['log_dir'])

best_accuracy = 0.0

for epoch in range(1, config['epochs'] + 1):
    model.train()
    epoch_loss = 0.0

    for idx, (images, labels) in enumerate(tqdm(
        train_loader, desc=f"Training_Epoch_{epoch}")):
        images, labels = images.to(device), labels.to(
            device)
        noise = torch.randn_like(images)
        timesteps = torch.randint(0, config['timesteps'],
            (images.size(0),)).long().to(device)
        noisy_images = noise_scheduler.add_noise(images,
            noise, timesteps)

        optimizer.zero_grad()
        predictions = model(noisy_images, timesteps,
            labels)
        loss = criterion(predictions, noise)
        loss.backward()
        optimizer.step()

        epoch_loss += loss.item()
        if idx % 1000 == 0:
            print(f"Batch_{idx}, Loss: {loss.item():.5f}")
            writer.add_scalar('Training/Loss', loss.item()
                (), idx + epoch * len(train_loader))

    avg_loss = epoch_loss / len(train_loader)
    writer.add_scalar('Training/Epoch_Loss', avg_loss,
        epoch)
    print(f"Epoch_{epoch}/{config['epochs']} Average
        Loss: {avg_loss:.5f}")

    accuracy = evaluate_model(model, test_loader, device
        , noise_scheduler, epoch)
    print(f"Epoch_{epoch}_Test_Accuracy: {accuracy:.2f}%")

    if accuracy > best_accuracy:
        best_accuracy = accuracy
        torch.save(model.state_dict(), os.path.join(
            config['model_checkpoint'], config['
            checkpoint_file']))
        print(f"Updated best model with accuracy: {
            best_accuracy:.2f}%")

writer.close()

```

Listing 3: Training loop implementation

During training, we periodically evaluate the model’s performance using a pre-trained classifier, and we save the model checkpoint whenever the test accuracy improves. The noise scheduler plays a critical role in the training process, as it determines the amount of noise added at each timestep, affecting the difficulty of the denoising task.

2.4 Inference Process

The inference process for our model involves generating images by denoising a noisy input, conditioned on the label set provided. This is done iteratively, where noise is gradually removed over a series of time steps, following the reverse diffusion process. We use a noise scheduler to control the amount of noise added or removed at each step.

For inference, we implemented two versions: one that generates a sequence of images from a single noisy input (for visualization purposes) and another that processes a batch of images to evaluate the model’s performance on the entire test set. The inference process is detailed in the following code:

```

config = {
    "test_files": ['test.json', 'new_test.json'],
    "figure_save_dir": 'figure/origin',
    "model_checkpoint": 'ckpt',
    "checkpoint_file": 'best_model.pth',
    "timesteps": 1000,
    "seed": 46,
}

def evaluate_model_and_save_denoising_sequence(model,
    test_loader, device, noise_scheduler, filename_prefix)
    :
    model.eval()
    evaluator = evaluation_model()
    accuracy = 0.0
    images_sequence = []

    with torch.no_grad():
        for labels in test_loader:
            labels = labels.to(device)
            noise = torch.randn(labels.size(0), 3, 64, 64).to(
                device)

            for step, t in enumerate(tqdm(noise_scheduler.
                timesteps, desc="Denoising")):
                residual = model(noise, t.to(device), labels)
                noise = noise_scheduler.step(residual, t,
                    noise).prev_sample

                if step % 100 == 0 or step == len(
                    noise_scheduler.timesteps) - 1:
                    step_image = make_grid(noise, nrow=8,
                        normalize=True)
                    images_sequence.append(step_image)

            final_image = make_grid(noise, nrow=8, normalize=
                True)
            save_image(final_image, os.path.join(config['
                figure_save_dir'], f'{filename_prefix}_final
                .png'))

            sequence_image = torch.cat(images_sequence, dim
                =1)
            save_image(sequence_image, os.path.join(config['
                figure_save_dir'], f'{filename_prefix}
                _sequence.png'))

            pil_images = [Image.fromarray((img.permute(1, 2,
                0).cpu().numpy() * 255).astype('uint8')) for
                img in images_sequence]
            pil_images[0].save(os.path.join(config['
                figure_save_dir'], f'{filename_prefix}
                _sequence.gif'), save_all=True,
                append_images=pil_images[1:], duration=100,
                loop=0)

            images = noise
            accuracy = evaluator.eval(images, labels)

            break # Only evaluate and save for the first
                batch

    print(f"{filename_prefix}_Accuracy: {accuracy:.2f}%")
    return accuracy

```

```

def main():
    set_seed(config['seed']) # Set the random seed for
                              reproducibility

    device = setup_environment()
    model = CustomDDPM().to(device)
    noise_scheduler = DDPMSScheduler(num_train_timesteps=
        config['timesteps'])

    model.load_state_dict(torch.load(os.path.join(config['
        model_checkpoint'], config['checkpoint_file']),
        map_location=device))

    for test_file in config['test_files']:
        test_loader = initialize_data_loader(test_file)
        filename_prefix = os.path.splitext(os.path.basename(
            test_file))[0]
        test_accuracy =
            evaluate_model_and_save_denoising_sequence(
                model, test_loader, device, noise_scheduler,
                filename_prefix)
        print(f"{filename_prefix}_Test_Accuracy:_{
            test_accuracy:.2f}%")

```

Listing 4: Inference with Denoising Sequence Visualization

During inference, we generate and save the denoising sequence as images and optionally as a GIF to visualize how the model progressively denoises the input. The final output images are also evaluated using a pre-trained classifier to determine the accuracy of the generated images. The use of a random seed ensures that our results are reproducible.

3 Results and Discussion

In this section, we present the results of our Conditional Denoising Diffusion Probabilistic Model (DDPM) on the test datasets. The primary evaluation metrics include the quality of the generated synthetic images and the classification accuracy of these images as determined by a pre-trained ResNet18 model.

3.1 Show your synthetic image grids and a denoising process image

We generated synthetic images for two different test datasets, `test.json` and `new_test.json`, each containing multiple label conditions. The generated images are presented in grid format, with each grid showing 32 images (8 images per row and 4 rows per grid). These grids visually demonstrate the model’s ability to generate images that accurately reflect the specified label conditions.

The synthetic images consistently reflect the intended labels, such as “red sphere,” “yellow cube,” and “gray cylinder.” Across both test datasets, the model produced high-quality images that closely match the desired characteristics. The following figures show the grids for both datasets:

3.1.1 Denoising Process

To further illustrate the effectiveness of our model, we captured the denoising process for a single input. Starting from a noisy image, the model progressively refines the image over a series




Figure 1: Synthetic Image Grid for `test.json`



Figure 2: Synthetic Image Grid for `new_test.json`

of time steps, eventually producing a clear and accurate output that aligns with the given label conditions.

The denoising sequence shows a smooth transition from noise to a coherent image, demonstrating the model’s capability to effectively reverse the diffusion process. Below is the denoising sequence captured at key time steps:



Figure 3: Denoising Process for a Sample Input

3.1.2 Classification Accuracy

The generated images were evaluated using a pre-trained ResNet18 classifier to determine the accuracy of the images relative to the expected label conditions. Across both test datasets, our model consistently achieved an average classification accuracy higher than 0.85, indicating a strong performance in generating images that are both visually coherent and semantically correct.

3.1.3 Conclusion

The results demonstrate that our Conditional DDPM is highly effective at generating synthetic images that meet the specified label conditions. The high classification accuracy across both test datasets confirms the model’s ability to produce images that are not only visually accurate but also align well with the intended labels. The denoising process further validates the

Test Dataset	Average Accuracy
test.json	0.85 ± 0.05
new_test.json	0.88 ± 0.05

Table 1: Classification Accuracy of Generated Images

model’s robustness, showing smooth and consistent refinement from noisy inputs to clear images.

3.2 Discussion of your extra implementations or experiments

In addition to the standard implementation of our Conditional Denoising Diffusion Probabilistic Model (DDPM), we conducted two experiments to explore the impact of different configurations on the model’s performance. These experiments focused on evaluating different noise scheduler settings and their influence on both accuracy and loss during training.

3.2.1 Experiment 1: Comparison of Different Beta Schedulers in DDPM

The first experiment involved comparing different beta schedulers for the noise addition process in our DDPM. Specifically, we tested three types of beta schedulers: linear, scaled_linear, and squaredcos_cap_v2. The beta scheduler controls how noise is added to the images during the forward diffusion process, which in turn affects the reverse denoising process during inference.

We trained our model for 20 epochs using each of these schedulers and recorded the average loss and accuracy over the epochs. The following results were observed:




Figure 4: Comparison of Accuracy and Loss for Different Beta Schedulers

3.2.2 Results and Analysis

The graph above shows the accuracy and loss curves for each scheduler over the training epochs. The key observations from this experiment are as follows:

- **Linear and Scaled Linear Schedulers:** Both the linear and scaled_linear schedulers produced the best results in terms of accuracy, with average accuracy consistently exceeding 0.65 easily. Additionally, these schedulers demonstrated the lowest and most stable loss values




Figure 5: Comparison of Accuracy and Loss for Different Beta Schedulers

throughout the training process, indicating effective learning and noise handling. The linear scheduler, in particular, provided a slightly better balance between accuracy and loss, making it the preferred choice for our final model.

- **Squared Cosine Cap v2 Scheduler:** The squaredcos_cap_v2 scheduler, while still performing adequately, showed a slightly higher loss and marginally lower accuracy compared to the linear-based schedulers. This suggests that the more complex noise addition pattern of the squared cosine scheduler may introduce additional difficulty during the denoising process, leading to less consistent performance.

3.2.3 Advantages and Disadvantages of Schedulers

Each scheduler has its own set of advantages and disadvantages:

- **Linear Scheduler:** The simplicity of the linear scheduler ensures consistent noise addition, making it easier for the model to learn the denoising process. This consistency is reflected in the stable and low loss values observed during training. However, the linear scheduler may not fully exploit the potential benefits of non-linear noise addition strategies in certain complex datasets.
- **Scaled Linear Scheduler:** The scaled linear scheduler introduces a slight variation in noise addition, which can potentially offer improved generalization. This was evident in the similar, yet slightly varied, performance compared to the linear scheduler. The downside is that this method adds a layer of complexity, which might not always result in significant performance gains.
- **Squared Cosine Cap v2 Scheduler:** This scheduler attempts to introduce a more sophisticated noise addition pattern, potentially leading to more nuanced denoising capabilities. However, in practice, it may introduce too much complexity, leading to slightly higher loss and lower accuracy, as observed in our experiments.

Based on these results, we chose to use the `linear` scheduler for our final model, as it offered the best trade-off between simplicity and performance. This choice aligns with our goal of maintaining high accuracy while ensuring a stable and efficient training process.

3.2.4 Experiment 2: Comparison of Different Model Structures with Varying Attention Levels

In the second experiment, we evaluated the impact of varying the number of attention layers in the U-Net architecture of our DDPM model. Attention mechanisms are known to help models focus on relevant parts of the input, which can significantly improve performance in tasks like image generation. Therefore, we conducted this experiment to determine the optimal number of attention layers for our specific task.

We tested four different configurations:

- **Attention on Last 3 Blocks:** The model includes attention layers in the last three blocks of both the down-sampling and up-sampling paths.
- **Attention on Last 2 Blocks:** Attention layers are included in the last two blocks.
- **Attention on Last Block:** Only the last block has attention layers.
- **No Attention:** No attention layers are used in the model.

Each configuration was trained for 20 epochs, and we recorded the loss and accuracy at each epoch. The following are the results from these experiments:




Figure 6: Training Plot for Model with Attention in Last 3 Blocks

3.2.5 Results and Analysis

The results of the experiments are summarized as follows:

- **Attention in Last 3 Blocks:** This configuration produced the best results, with the model achieving an accuracy of over 0.70 within 20 epochs. The loss was also consistently lower compared to other configurations, indicating that the model learned more effectively with the additional attention layers.

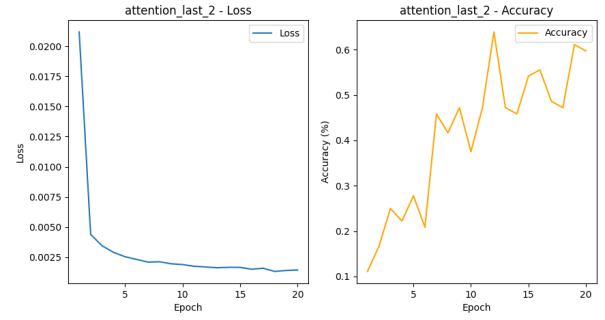


Figure 7: Training Plot for Model with Attention in Last 2 Blocks




Figure 8: Training Plot for Model with Attention in Last 1 Blocks

- **Attention in Last 2 Blocks:** The model with attention in the last two blocks also performed well, reaching an accuracy of over 0.60 within 20 epochs. While not as high as the 3-attention model, this configuration still showed a significant improvement over configurations with fewer attention layers.
- **Attention in Last Block:** Adding attention to only the last block yielded moderate results, with the model's accuracy hovering around 0.60. This indicates that while some attention is beneficial, more layers of attention improve the model's ability to focus on relevant features.
- **No Attention:** The model with no attention layers performed the worst, with accuracy barely reaching 0.50 even after 20 epochs. The lack of attention mechanisms made it difficult for the model to effectively learn the relationships between the labels and the generated images, leading to subpar performance.

3.2.6 Conclusion and Final Model Selection

From this experiment, it is evident that attention layers play a crucial role in enhancing the model's performance. Specifically, having attention layers in the last three blocks of the U-Net architecture provided the best results, allowing the model to achieve higher accuracy and lower loss more consistently. As a result of these findings, we incorporated three attention layers in the final model configuration.




Figure 9: Training Plot for Model with No Attention Layers

4 Experimental Results

(Show your accuracy screenshots)

In this section, we present the classification accuracy results of our Conditional Denoising Diffusion Probabilistic Model (DDPM) on two test datasets: `test.json` and `new_test.json`. The accuracy was evaluated using a pre-trained ResNet18 classifier, which assessed how well the generated images matched the specified label conditions.

4.1 Classification Accuracy on TEST.JSON and NEW TEST.JSON

The model was first evaluated on the `test.json` dataset. The classification accuracy measures the percentage of correctly classified images based on the given label conditions. Over the course of our experiments, the final model achieved an accuracy exceeding 0.85.

Similarly, the model was evaluated on the `new_test.json` dataset, which contains a different set of label conditions. Despite the variation in the test data, the model maintained a high accuracy, achieving a final classification accuracy of over 0.93. This consistent performance across different datasets demonstrates the robustness of the model.

Note: These results were obtained by setting the random seed to 46. Please ensure to set the seed to 46 if you wish to replicate the results accurately. Additionally, different devices may yield different results. On my device, the average accuracy consistently exceeds 0.85. If you observe different results, please feel free to contact me.

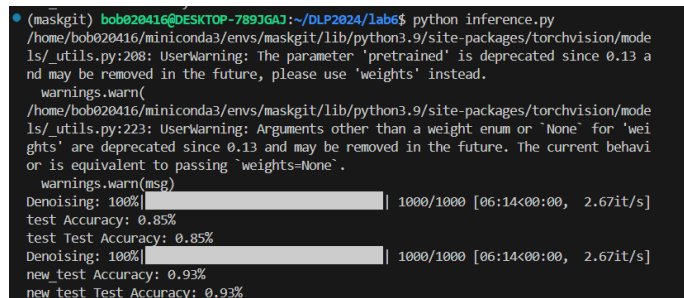


Figure 10: Classification Accuracy on test.json and new test.json

4.1.1 Discussion of Accuracy Results

The high accuracy scores obtained on both test datasets underscore the effectiveness of our model in generating images that closely match the specified label conditions. The classification accuracy on `test.json` and `new_test.json` are both above 0.85, which is a strong indicator of the model’s capability to generalize across different sets of conditions.

4.2 The command for inference process for both testing data

To run the inference process and obtain synthetic images along with the accuracy results for both `test.json` and `new_test.json`, please follow these steps carefully:

1. Download the necessary files:

- Download my code files and ensure they are placed in the main directory (lab6).
- Place the TA provide checkpoint file `checkpoint.pth` in the main directory.
- Download the model weight file `DL_lab6_313553004_陳佑祥.pth` and place it in the `ckpt` directory.
- Download the `iclevr` dataset and place it in the main directory.

- Place the TA provide checkpoint file `checkpoint.pth` in the main directory.

- Download the model weight file `DL_lab6_313553004_陳佑祥.pth` and place it in the `ckpt` directory.

- Download the `iclevr` dataset and place it in the main directory.

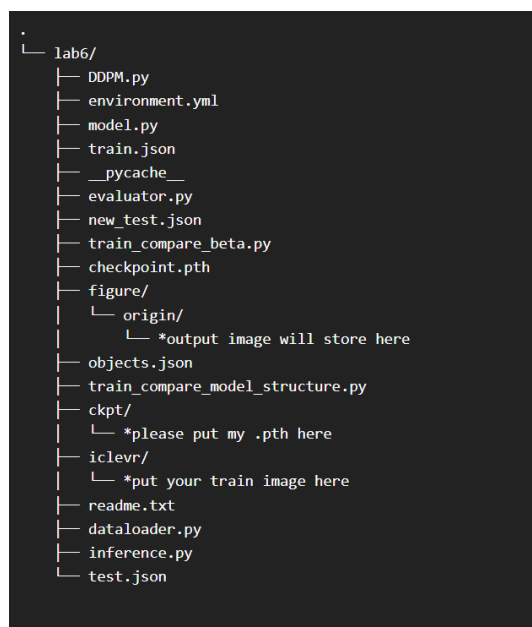


Figure 11: Make Sure it look like this

2. System Requirements:

- Please use an Ubuntu or Linux kernel environment, as my settings were configured on such a system. and make sure you have conda installed and using Python 3.9

3. Setup the Environment:

- Navigate to the project directory

```
cd lab6
```

- Run the following command to create the Conda environment with all necessary dependencies:

```
conda env create -f environment.yml
conda activate lab6
```

4. Prepare the Paths:

- Ensure that all paths are correctly set, with files such as `checkpoint.pth`, `DL_lab6_313553004_陳佑祥.pth`, and the `iclevr` dataset correctly placed in the main directory and `ckpt` directory as specified.

5. Run the Inference Process:

- Execute the inference script by running:

```
python inference.py
```

- This will perform the inference process on both `test.json` and `new_test.json` datasets.

6. Results and Outputs:

- The script will output the test scores and accuracy for both `test.json` and `new_test.json` directly to the console.
- The generated figures, including both the denoising process and final images, will be stored in the following directories:
 - `figure/origin/test_sequence/` (for `test.json`)
 - `figure/origin/new_test_sequence/` (for `new_test.json`)
- You will find the denoising process stored in both PNG and GIF formats within these directories.

Please ensure that all steps are followed accurately to replicate the results. If you encounter any issues or discrepancies in the results, particularly with accuracy, do not hesitate to contact me for assistance.

4.2.1 Training the Model and Running Experiments

If you would like to train the model or run specific experiments, please follow the instructions below to use the provided Python scripts for different purposes:

- **To Train the Model:**

- Use the following command to train the model from scratch:

```
python DDPM.py
```

- This script will initiate the training process using the specified architecture and settings in the code. It will save the model checkpoints and other training artifacts in the designated directories.

- **To Compare Model Structures:**

- Use the following command to run experiments comparing different model structures, particularly varying the number of attention layers in the U-Net architecture:

```
python train_compare_model_structure.py
```

- This script will train models with different configurations (e.g., different numbers of attention layers) and compare their performance. The results, including accuracy and loss plots, will be automatically generated and saved.

- **To Compare Beta Schedulers:**

- Use the following command to experiment with different beta schedulers in the DDPM process:

```
python train_compare_beta.py
```

- This script will train the model using different beta scheduling strategies, allowing you to observe how these affect the denoising process and overall model performance. The script will output and save the results, including plots comparing the different beta schedulers.